

COBY KASSNER

coby.lk | [linkedin.com/in/cobyk/](https://www.linkedin.com/in/cobyk/)
ansatz@coby.lk | +1 (720) 551-3481

Technically-oriented AI safety generalist with interest in automated safety, generalization, and macrostrategy.

RESEARCH EXPERIENCE

Interpretability by design | [main](#) + [extension](#)

February–August 2025

Supervised Program for Alignment Research (SPAR)

- Designed and tested simplex-constrained MLPs towards inherently interpretable models, mentored by Ronak Mehta
- In extension of SPAR, developed a key residual-stream constraint technique that enabled stable training to low loss, after several prior simplex-constrained transformer attempts failed
- Implemented attention and feed-forward layers in centered log-ratio coordinates, with residual updates mapped back via simplex renormalization so each layer state remained a valid distribution

Inference-time privacy editing for LLMs | [link](#)

September–December 2025

Independent

- Borrowed activation editing techniques from mechanistic interpretability research and applied them to privacy, using internal interventions to suppress leaked personal information at inference time
- Reduced targeted leakage by up to 65.6% on GPT-Neo-1.3B and found a simple intervention could match a more complex one, supporting lower-cost deployment

OPERATIONS EXPERIENCE

Executive Board

September 2025–Present

Yale AI Alignment (YAIA), *Yale Effective Altruism (YEA)*

- Conducted end-to-end operations for the largest-ever iterations of both organizations' flagship fellowship programs, leading program strategy, recruiting, and logistics
- Managed the full application processes and built a custom platform for forms and application management, with pipelines for applications and scheduling that enabled accepting 125 fellows for YEA's summer fellowship
- Rebuilt YAIA's fellowship curriculum from scratch and revamped organization-wide branding

Developer, Tracks Program

June 2026–Present

Existential Risk Laboratory at the University of Chicago (XLab)

- Building interactive, mnemonic courses for AI safety, designing materials, demos, and exercises
- Helping structure pedagogical strategy to optimize for learning outcomes

Staff Team

2023–2024

International Research Olympiad (IRO)

- Directed program to start over 320 research clubs in secondary schools across 40 countries
- Collaborated with leadership team to coordinate over 50 student volunteers, negotiate over \$15,000 in sponsorships, and organize in-person finals event in Cambridge, MA

TECHNICAL SKILLS

Research: LLM fine-tuning and inference (LoRA, open-weight models to 70B), mechanistic interpretability, activation steering/editing, custom transformer architectures, alignment evaluation and red-teaming, privacy techniques

Libraries: PyTorch, Transformers (HuggingFace), Transformer Lens, NumPy, Pandas

Languages: Python (6 years), C++ (3 years)

EDUCATION

Statistics and Data Science, B.S.

2025–2029

Yale College

Coursework: Trustworthy Deep Learning, Large Language Models, Intermediate ML, Probability Theory, Theory of Statistics, Intensive Linear Algebra, Intensive Analysis, Superintelligence (Writing Seminar)

Computer Science, A.S. and Mathematics, A.S.

2021–2025

Arapahoe Community College (concurrent enrollment)

Coursework: Linear Algebra, Multivariate Calculus, Data Structures & Algorithms, Computer Architecture

High School Diploma

2021–2025

Colorado Early Colleges Douglas County North

Class rank: 4/209