

COBY L. KASSNER

cobylk.io | [linkedin.com/in/cobylk](https://www.linkedin.com/in/cobylk)
kassner@cobylk.io | +1 (720) 551-3481

Researcher with interest in interpretability, model psychology, and more broadly, applied AI safety.

SELECTED RESEARCH

Exploring Emergent Misalignment | *(In progress)*

February 2026–Present

- Attempting to characterize what features of a dataset cause it to elicit emergent misalignment and exploring several methods for predicting whether or not a given dataset will elicit unsafe behavior
- Ran many (>50) fine-tuning runs on open source models up to the 70B scale, conducted inference and various evaluations on these models, and did some mechanistic work (e.g., steering using a rank-1 LoRA decoder)

Inference-Time Privacy Editing for LLMs | [link](#)

September–December 2025

- Borrowed activation editing techniques from mechanistic interpretability research and applied them to privacy, using internal interventions to suppress leaked personal information at inference time
- Reduced targeted leakage by up to 65.6% on GPT-Neo-1.3B and found a simple intervention could match a more complex one, supporting lower-cost deployment

Interpretable-by-design Transformers | [link](#)

May–August 2025

- In extension of SPAR work, developed a key residual-stream constraint technique that enabled stable training to low loss, after several prior simplex-constrained transformer attempts failed
- Implemented attention and feed-forward layers in centered log-ratio coordinates, with residual updates mapped back via simplex renormalization so each layer state remained a valid distribution

EXPERIENCE

Research Fellow

February–May 2025

Supervised Program for Alignment Research ([SPAR](#))

- Conducted [interpretability research](#) under mentorship from Dr. Ronak Mehta
- Designed and tested simplex-constrained MLPs towards inherently interpretable models

Staff Team

2023–2024

International Research Olympiad ([IRO](#))

- Directed program to start over 320 research clubs in secondary schools across 40 countries
- Collaborated with leadership team to coordinate over 50 student volunteers, negotiate over \$15,000 in sponsorships, and organize in-person finals event in Cambridge, MA

TECHNICAL SKILLS

Research: Mechanistic interpretability, activation steering/editing, LLM fine-tuning and inference, custom transformer architectures, privacy techniques, alignment evaluation and red-teaming

Libraries: Transformers (HuggingFace), PyTorch, JAX, NumPy, Pandas, Scikit-Learn, Transformer Lens

Languages: Python (6 years), C++ (3 years)

EDUCATION

Statistics and Data Science, B.S.

2025–2029

Yale College

Activities: AI Alignment (Board, Head of Operations), Effective Altruism (Board, Intro Fellowship Manager)

Coursework: Trustworthy Deep Learning, Large Language Models, Intermediate ML, Probability Theory, Theory of Statistics, Intensive Linear Algebra, Intensive Analysis, Superintelligence (Writing Seminar)

Computer Science, A.S. and Mathematics, A.S.

2021–2025

Arapahoe Community College (concurrent enrollment)

Coursework: Linear Algebra, Multivariate Calculus, Data Structures & Algorithms, Comp. Architecture

High School Diploma

2021–2025

Colorado Early Colleges Douglas County North

Class rank: 4/209