

COBY KASSNER

coby.lk | [linkedin.com/in/cobyk/](https://www.linkedin.com/in/cobyk/)
ansatz@coby.lk | +1 (720) 551-3481

Technically-oriented AI safety generalist with interest in
macrostrategy, control, and automated safety research.

RESEARCH EXPERIENCE

Directions in risk-averse AI August 2026–Present

Ongoing independent technical project mentored by Pablo Bernabeu Pérez

Topics in strategy and futurism August 2026–Present

Ongoing SPAR project mentored by Dylan Bowman, Govind Pimpale, and Blaine Dillingham

Interpretability by design | [main](#) + [extension](#) February–August 2025

Supervised Program for Alignment Research (SPAR)

- Designed and tested simplex-constrained MLPs towards inherently interpretable models, mentored by Ronak Mehta
- In extension of SPAR, developed a key residual-stream constraint technique that enabled stable training to low loss, after several prior simplex-constrained transformer attempts failed
- Implemented attention and feed-forward layers in centered log-ratio coordinates, with residual updates mapped back via simplex renormalization so each layer state remained a valid distribution

Inference-time privacy editing for LLMs | [link](#) September–December 2025

Independent

- Borrowed activation editing techniques from mechanistic interpretability research and applied them to privacy, using internal interventions to suppress leaked personal information at inference time

GENERALIST EXPERIENCE

Executive Board September 2025–Present

President of Yale AI Alignment (YAIA), Fellowship Coordinator for Yale Effective Altruism (YEA)

- Conducted end-to-end operations for the largest-ever iterations of both organizations' flagship fellowship programs, leading program strategy, recruiting, and logistics
- Managed the full application processes and built a custom platform for forms and application management, with pipelines for applications and scheduling that enabled accepting 125 fellows for YEA's summer fellowship
- Rebuilt YAIA's fellowship curriculum from scratch and revamped organization-wide branding

Developer, Tracks Program June 2026–Present

Existential Risk Laboratory at the University of Chicago (XLab)

- Building interactive, mnemonic courses for AI safety, designing materials, demos, and exercises
- Helping structure pedagogical strategy to optimize for learning outcomes

Staff Team 2023–2024

International Research Olympiad (IRO)

- Directed program to start over 320 research clubs in secondary schools across 40 countries
- Collaborated with leadership team to coordinate over 50 student volunteers, negotiate over \$15,000 in sponsorships, and organize in-person finals event in Cambridge, MA

TECHNICAL SKILLS

Research: LLM fine-tuning and inference (LoRA, open-weight models to 70B), mechanistic interpretability, activation steering/editing, custom transformer architectures, alignment evaluation and red-teaming, privacy techniques

Libraries: PyTorch, Transformers (HuggingFace), Transformer Lens, NumPy, Pandas

Languages: Python (6 years), C++ (3 years)

EDUCATION

Statistics and Data Science, B.S. 2025–2029

Yale College

Coursework: Frontier AI, Trustworthy Deep Learning, Large Language Models, Intermediate ML, Probability Theory, Theory of Statistics, Intensive Linear Algebra, Intensive Analysis, Superintelligence (Writing Seminar)

Computer Science, A.S. and Mathematics, A.S. 2021–2025

Arapahoe Community College (concurrent enrollment)

Coursework: Linear Algebra, Multivariate Calculus, Data Structures & Algorithms, Computer Architecture